Computer Networks
Programming Assignment

Team Members

Alaa Mohamed Abdel Hamid	**6473**
Nouran Hisham	**6532**
Kareem Sabra	**6491**

Multi-threaded web server:

In this part we used multi-threading and socket programming concepts. Our server is connected on port 8085 and gives out an error if this port is not free to connect to.

The server makes TCP connections with its clients and responds to requests of type GET or POST only with either 200 OK or 404 not found.

Our server responds to different HTTP versions, HTTP 1.0 and HTTP 1.1. The main differences between the two versions are persistent connections and pipelining which are present in HTTP 1.1 and not present in HTTP 1.0.

In HTTP 1.0, we immediately close the client socket after sending the response to the client. However, in Http 1.1 we keep the socket opened for 10 seconds ready to receive any further requests from the same client without having to connect to our server again.

All our server files in a folder called template which the server searches in for any files requested by the client in GET requests.

Server sample runs:

This is when we first run our server:

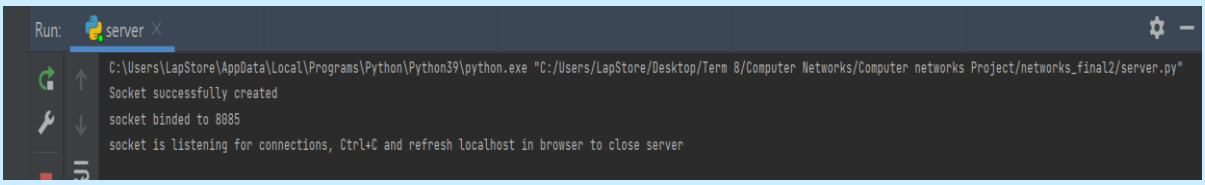

This is testing our server with postman:
GET request (html file):

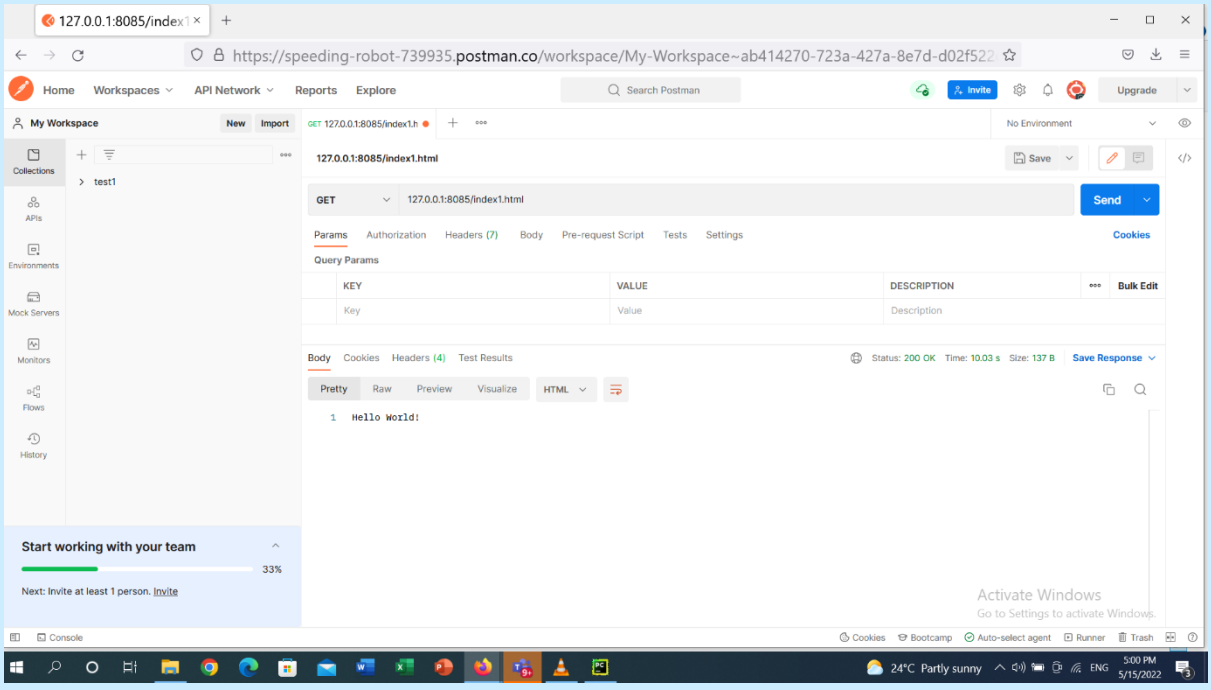

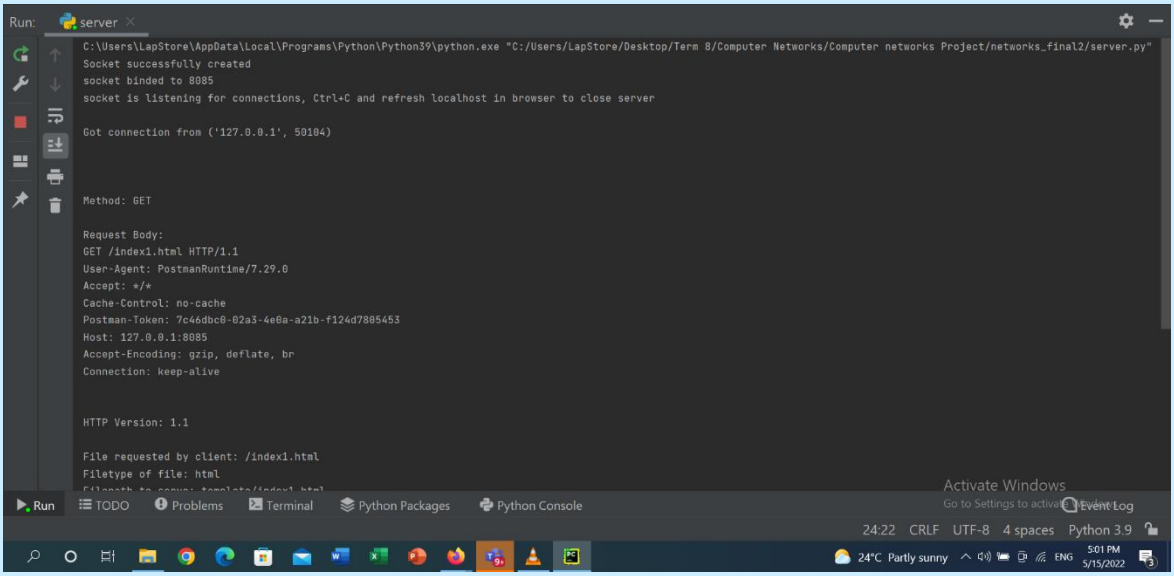

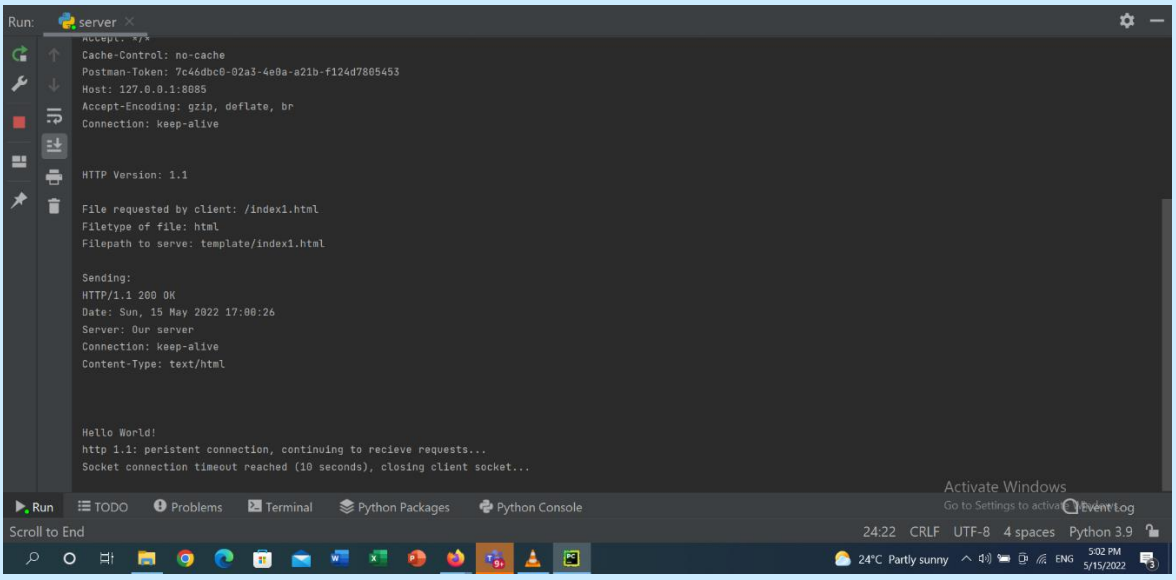

GET request (png file):

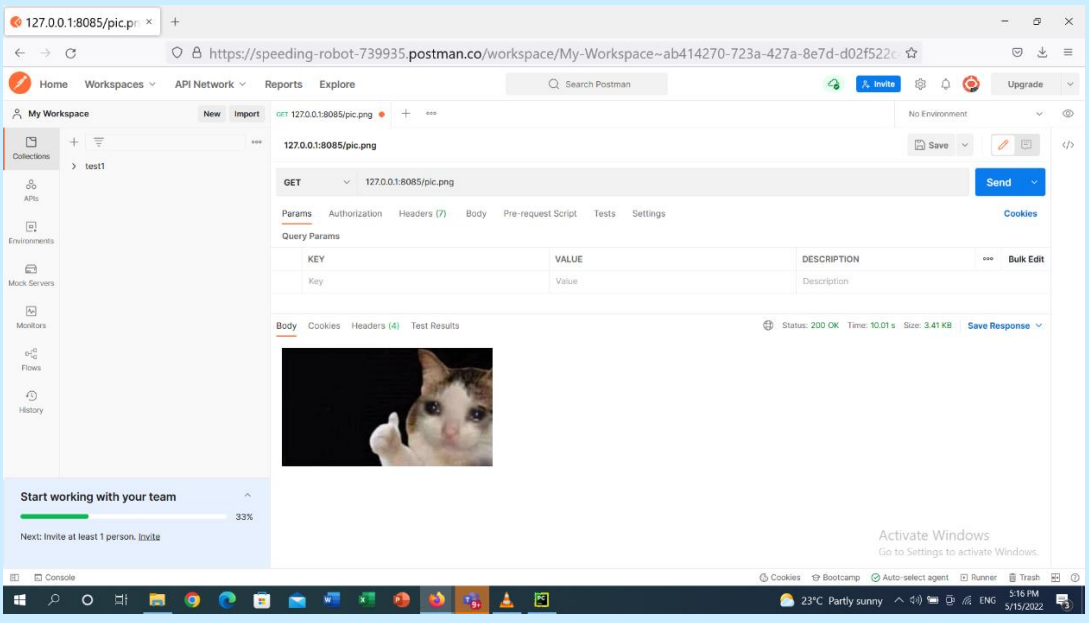

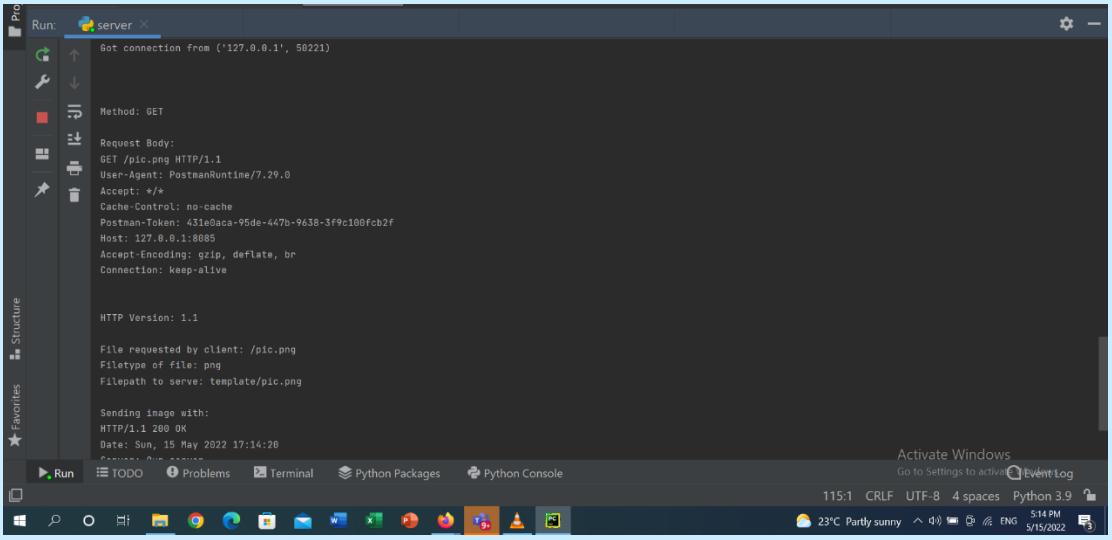

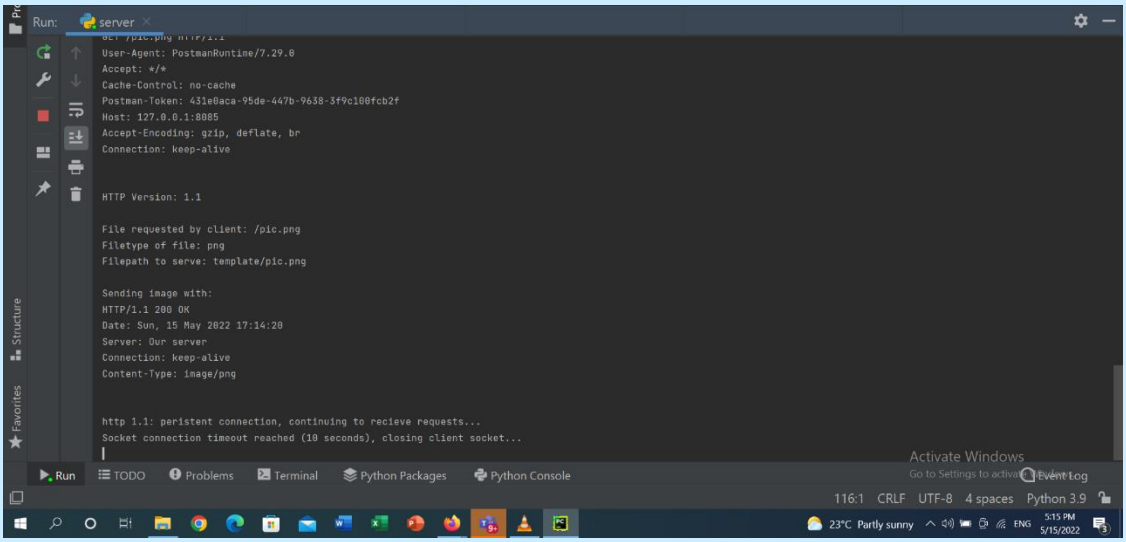

POST request (txt file):

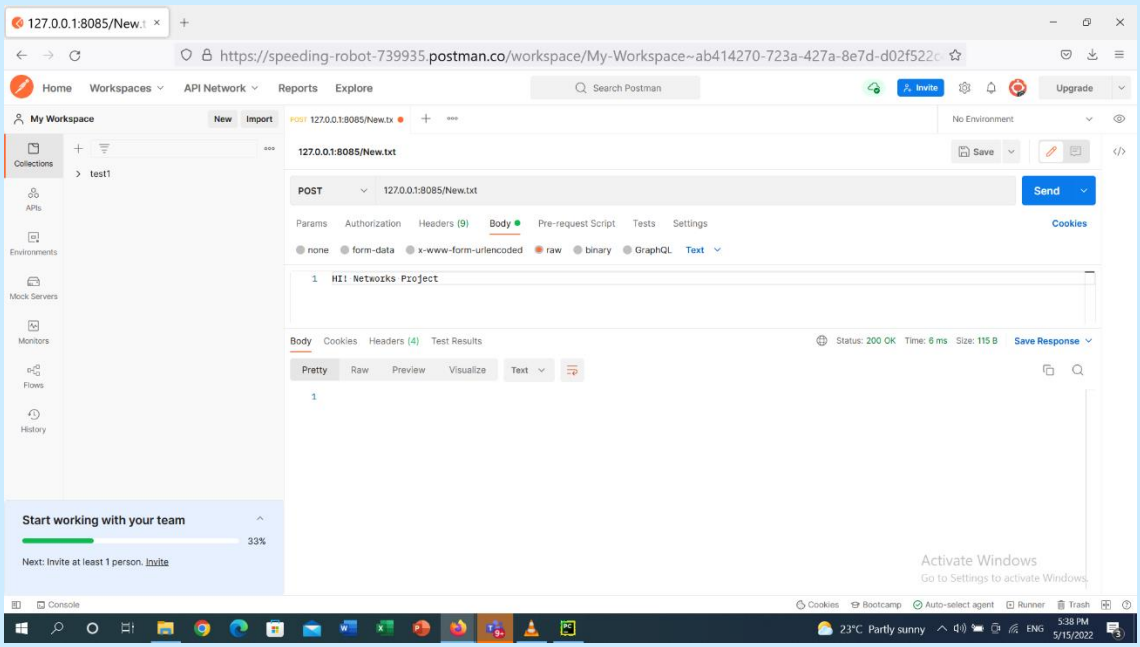

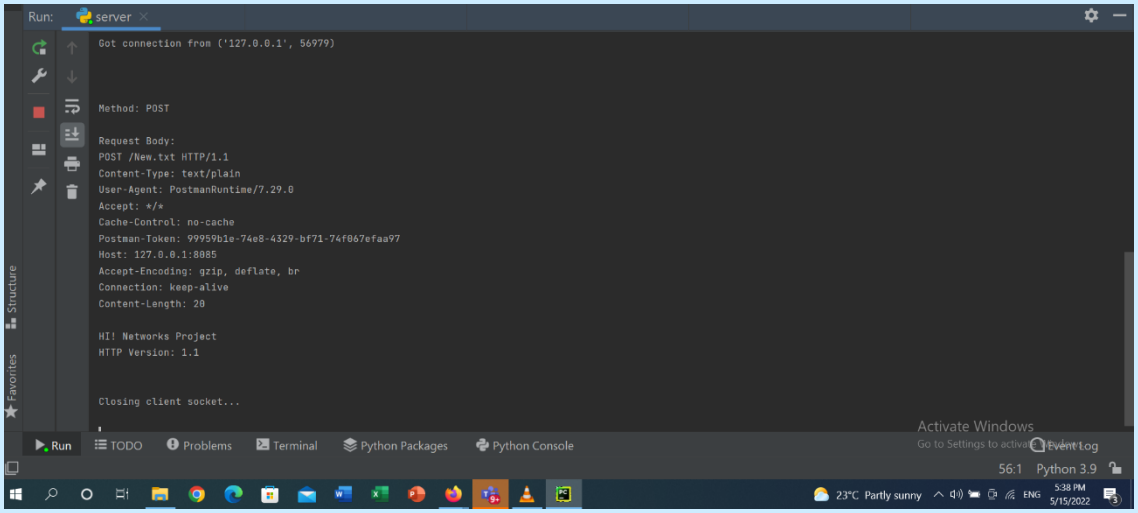

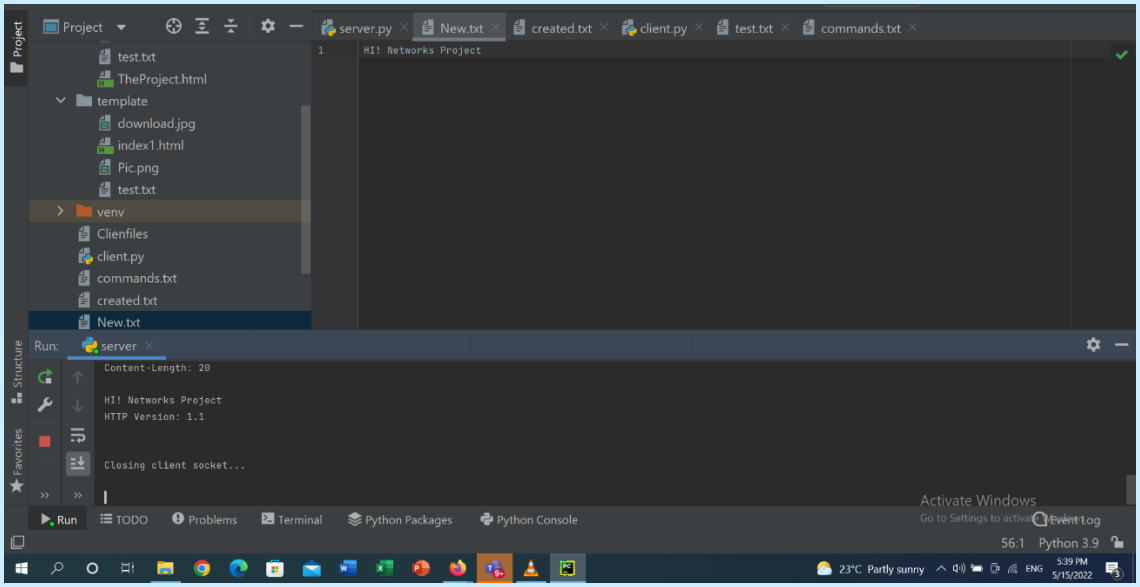

HTTP web client:

We made our web client using selectors, the client reads a (commands.txt) file in which commands are written in the following format (GET/POST file-name host-name {port-number}). These commands are then parsed to generate the HTTP requests accordingly.

Once the request is created it's sent to our start_connection function to initiate connection with the wanted host.

We tested our client on both our server and info.cer.ch as well as wireshark test.

Client sample runs:

Wireshark testing:

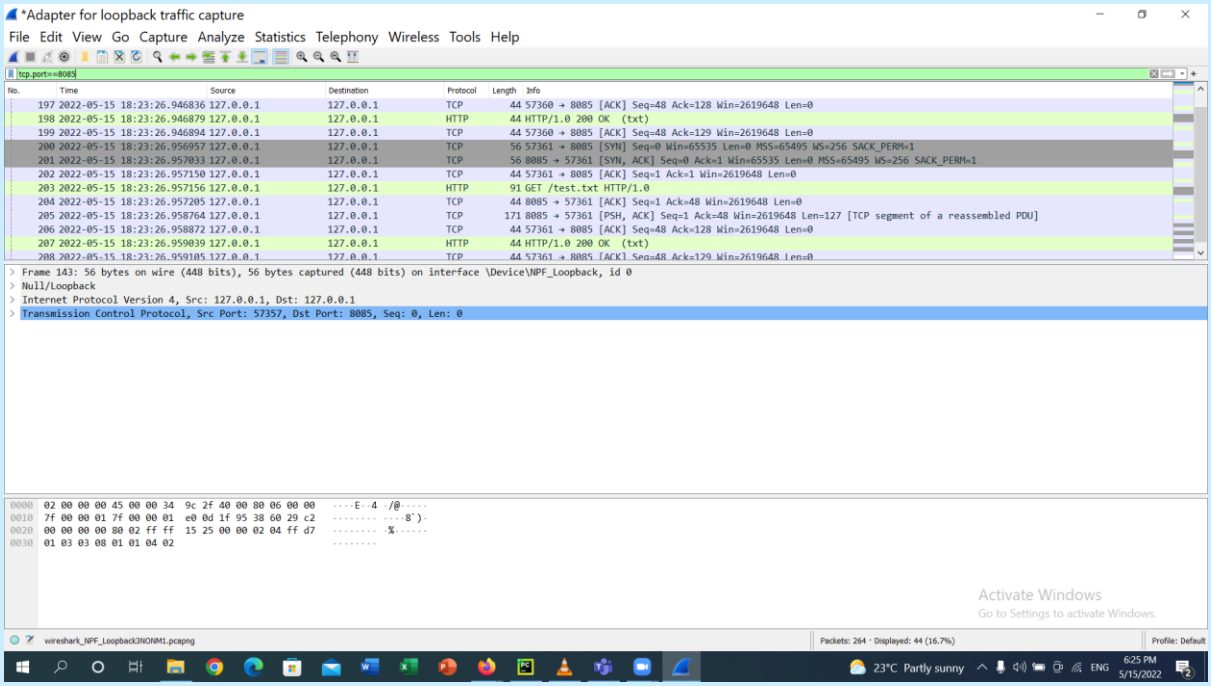

Testing our client with info.cern.ch:

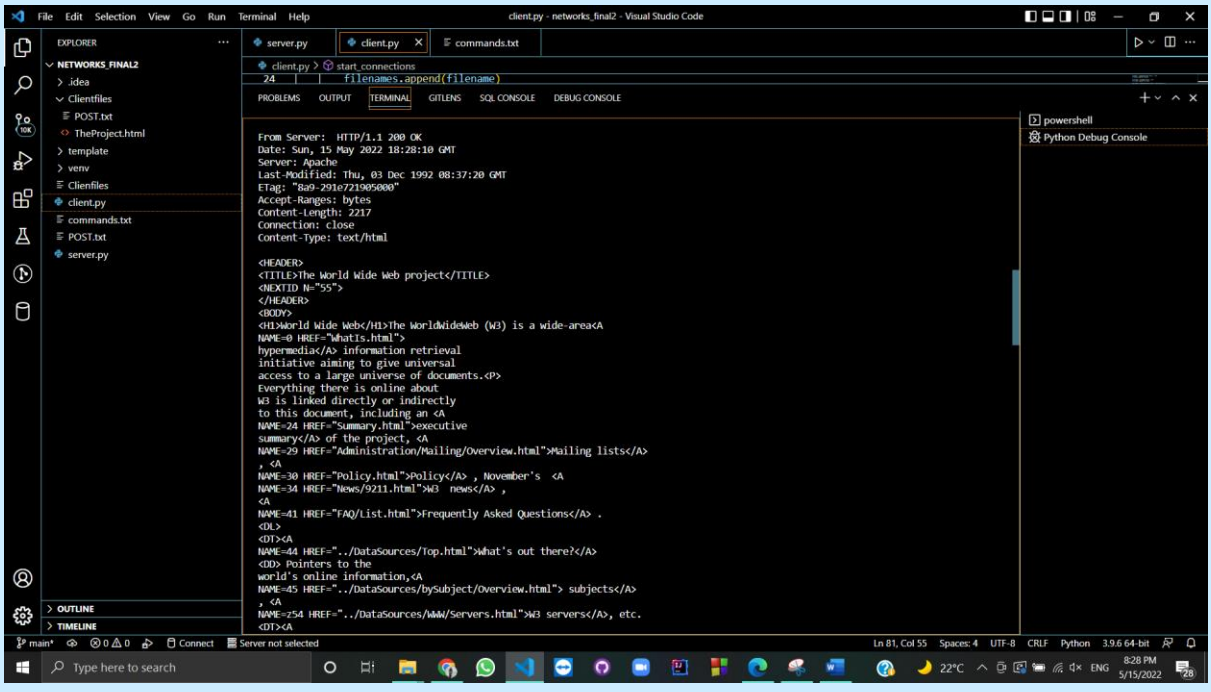

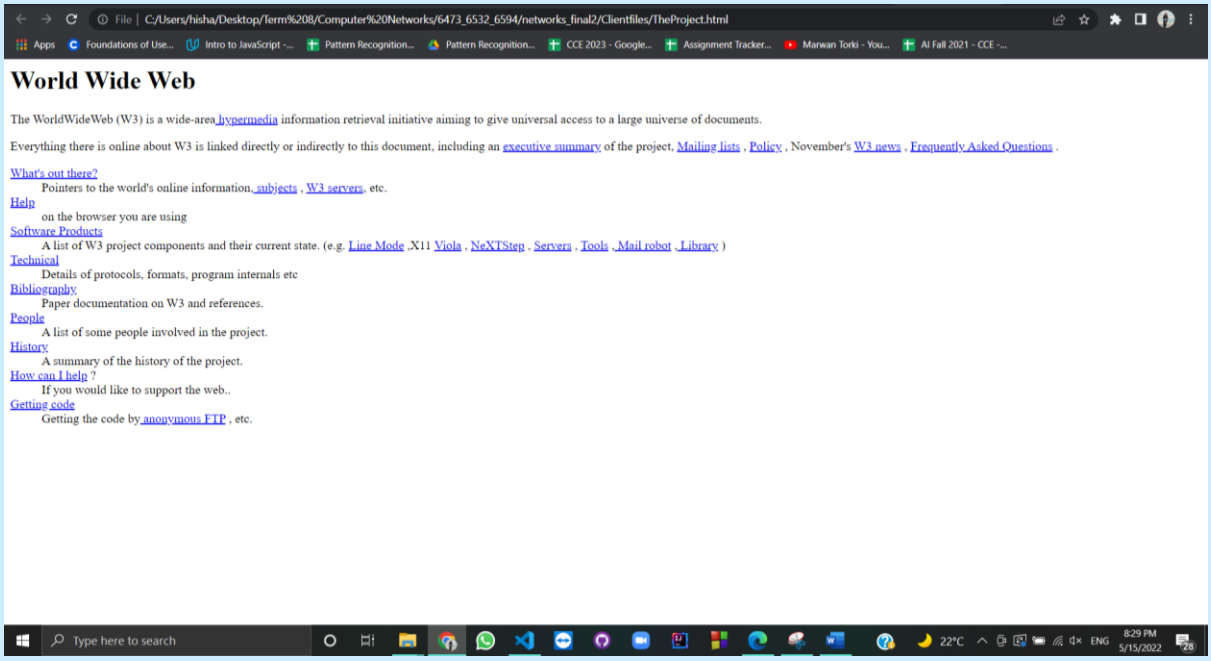

We couldn't test the POST requests of our client on any known server because simply, we wouldn't be able to see that the file was actually created in the server, so we tested this part on our own server.

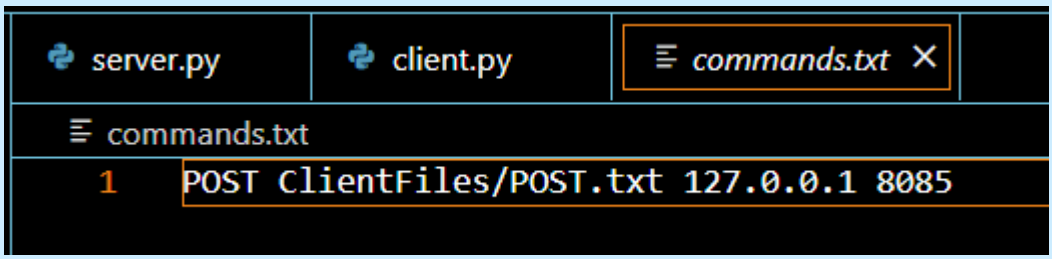

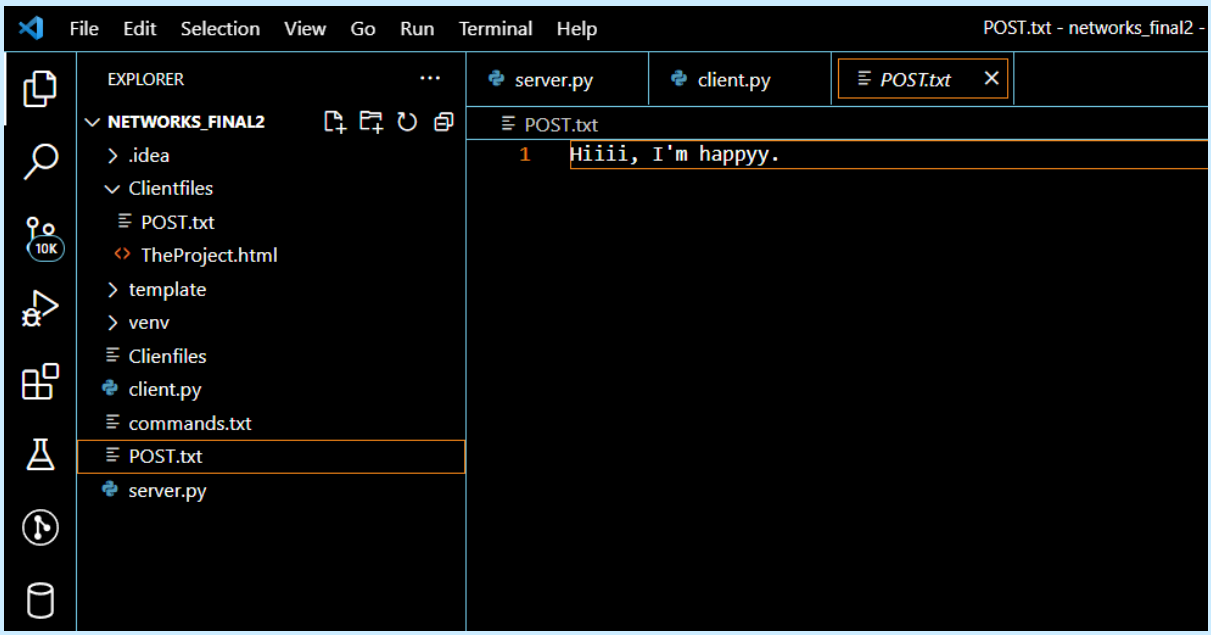

Bonus:

Additionally, we implemented a simple cache in our client to prevent uncalled for connections when the same request has been made before.

We just saw how we connected our client to info.cern.ch, we'll try to send the same exact request and we'll see that our client gets the response this time from the cache without connecting to the server again.